

CS203: Software Tools & Techniques for AI

Lab 01: Distributed Tracing and Telemetry (https://github.com/nishchaybhutoria/CS203_Lab_01-main)

IIT Gandhinagar, Sem II - 2024-25

Team Members

- Nishchay Bhutoria (23110222)
 - Srivaths P (23110321)
-

Introduction

This lab submission demonstrates the use of distributed tracing and telemetry in a Flask-based Course Information Portal. Using **OpenTelemetry** and **Jaeger**, we add observability to the application, enabling efficient debugging, performance monitoring, and error tracking.

Technologies Used

- **Flask**: Web framework for building the portal.
 - **OpenTelemetry**: Framework for generating and exporting telemetry data.
 - **Jaeger**: Distributed tracing backend for storing and visualizing traces.
 - **python-json-logger**: For structured JSON logging.
-

Features

1. Add Courses to the Catalog

- An "Add a New Course" button on the catalog page navigates to a form for adding courses.

Course Catalog

- [Introduction to Computer Science](#) - CS101
- [Software and Tools for AI](#) - CS 203
- [Data Structures and Algorithms II](#) - ES 301

Add a course

Add a Course

Course Code

Course Name

Instructor

Semester

Schedule

Classroom

Prerequisites

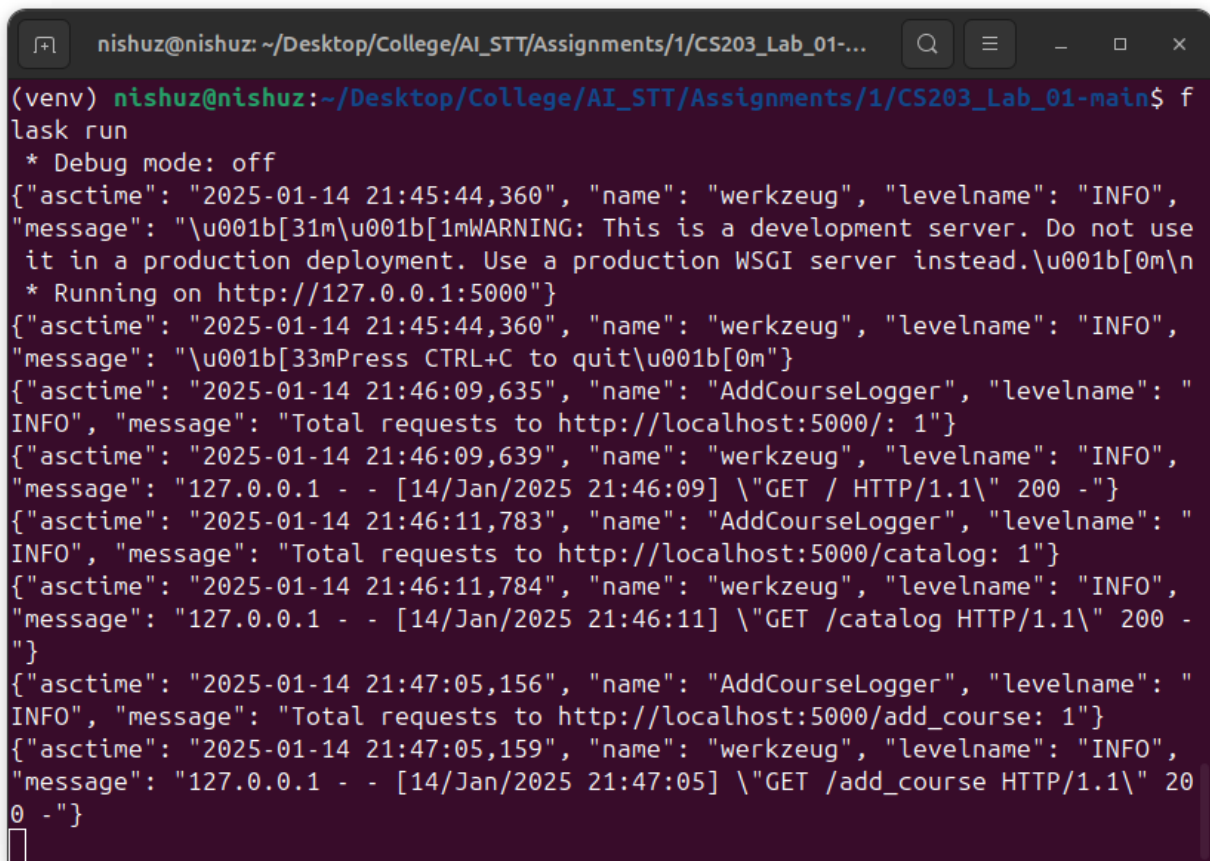
Grading

Description

- Logs are generated for:
 - **Successful course additions** (**INFO** level).
 - **Missing required fields** (**ERROR** level).
- Example log:

```
{
  "asctime": "2025-01-14 19:21:42,571",
  "name": "AddCourseLogger",
  "levelname": "INFO",
  "message": "Course CS 201 added successfully. All required fields
are present."
}
```

- Console logs:

A terminal window with a dark background and light-colored text. The prompt is 'nishuz@nishuz: ~/Desktop/College/AI_STT/Assignments/1/CS203_Lab_01-...'. The user has entered 'f' and 'lask run'. The output shows Flask's debug mode is off, the server is running on http://127.0.0.1:5000, and several log messages from Werkzeug and AddCourseLogger. The logs show requests to /, /catalog, and /add_course, all returning 200 status codes.

```
(venv) nishuz@nishuz: ~/Desktop/College/AI_STT/Assignments/1/CS203_Lab_01-...$ f
lask run
* Debug mode: off
{"asctime": "2025-01-14 21:45:44,360", "name": "werkzeug", "levelname": "INFO",
"message": "\u001b[31m\u001b[1mWARNING: This is a development server. Do not use
it in a production deployment. Use a production WSGI server instead.\u001b[0m\n
* Running on http://127.0.0.1:5000"}
{"asctime": "2025-01-14 21:45:44,360", "name": "werkzeug", "levelname": "INFO",
"message": "\u001b[33mPress CTRL+C to quit\u001b[0m"}
{"asctime": "2025-01-14 21:46:09,635", "name": "AddCourseLogger", "levelname": "
INFO", "message": "Total requests to http://localhost:5000/: 1"}
{"asctime": "2025-01-14 21:46:09,639", "name": "werkzeug", "levelname": "INFO",
"message": "127.0.0.1 - - [14/Jan/2025 21:46:09] \"GET / HTTP/1.1\" 200 -"}
{"asctime": "2025-01-14 21:46:11,783", "name": "AddCourseLogger", "levelname": "
INFO", "message": "Total requests to http://localhost:5000/catalog: 1"}
{"asctime": "2025-01-14 21:46:11,784", "name": "werkzeug", "levelname": "INFO",
"message": "127.0.0.1 - - [14/Jan/2025 21:46:11] \"GET /catalog HTTP/1.1\" 200 -
"}
{"asctime": "2025-01-14 21:47:05,156", "name": "AddCourseLogger", "levelname": "
INFO", "message": "Total requests to http://localhost:5000/add_course: 1"}
{"asctime": "2025-01-14 21:47:05,159", "name": "werkzeug", "levelname": "INFO",
"message": "127.0.0.1 - - [14/Jan/2025 21:47:05] \"GET /add_course HTTP/1.1\" 20
0 -"}

```

2. OpenTelemetry Tracing

We instrumented the following routes to provide detailed traces for observability:

Route: `/catalog`

- **Span Name:** `render_catalog`
- **Attributes:**
 - `request_method`: HTTP method (e.g., `GET`).
 - `user_ip`: User's IP address.
 - `course_count`: Total number of courses loaded.

Route: `/add_course`

1. **Span Name:** `view_add_course_form` (GET request)
 - **Description:** Tracks the rendering of the course addition form.
 - **Attributes:**
 - `request_method`: HTTP method (e.g., `GET`).
 - `user_ip`: User's IP address.
2. **Span Name:** `submit_add_course_form` (POST request)
 - **Description:** Tracks the submission of the course addition form.

- **Attributes:**

- **request_method**: HTTP method (e.g., **POST**).
- **user_ip**: User's IP address.
- **course_code**: Code of the course being added.
- **course_name**: Name of the course being added.

3. Span Name: **add_course_form_validation_error** (Child Span)

- **Description**: Captures validation errors during form submission.

- **Attributes:**

- **missing_fields**: Fields that are missing in the form.
- **error_count**: Number of errors (client-side or server-side).

4. Span Name: **count_errors** (Child Span)

- **Description**: Tracks the total number of client-side and server-side errors.

- **Attributes:**

- **client_error_count**: Count of client-side errors.
- **server_error_count**: Count of server-side errors.

Route: **/course/<code>**

- **Span Name**: **view_course**

- **Attributes:**

- **request_method**: HTTP method (e.g., **GET**).
- **user_ip**: User's IP address.
- **course_code**: Code of the course being viewed.
- **error**: Set to **True** if the course is not found.
- **error_message**: Error message if the course is not found.

3. Exporting Telemetry Data to Jaeger

- Traces include:
 - Total requests to each route.
 - Processing time for operations.
 - Error counts for client-side and server-side errors.
- Screenshot of Jaeger traces:

Missing Fields:

Submit

- "instructor" field is missing
- "semester" field is missing
- "schedule" field is missing
- "classroom" field is missing
- "prerequisites" field is missing
- "grading" field is missing

JAEGER UI

Search

Compare

System Architecture

Monitor

catalog-app: POST /add_course 90ab084

Find...

Trace Timeline

Trace Start January 14 2025, 21:00:19.739 Duration 1.74ms Services 1 Depth 3 Total Spans 5

Service & Operation

catalog-app POST /add_course

catalog-app count_requests

catalog-app submit_add_course_form

catalog-app add_course_form_validation_error

catalog-app count_errors

add_course_form_validation_error

count_errors

The trace data contains exactly which fields are missing and also has an explicit error type.

Successful Course Addition:

JAEGER UI

Search

Compare

System Architecture

Monitor

catalog-app: POST /add_course 7c5fab3

Find...

Trace Timeline

Trace Start January 14 2025, 21:03:25.829 Duration 1.55ms Services 1 Depth 3 Total Spans 3

Service & Operation

catalog-app POST /add_course

catalog-app count_requests

catalog-app submit_add_course_form

POST /add_course

count_requests

submit_add_course_form

/

As visible from the above image, we are able to get metadata about the form submission, such as the course code, course name, etc. as well.

4. JSON Logging

- Structured logs for all events:
 - Example:

```
{
  "asctime": "2025-01-14 19:21:07,417",
  "name": "AddCourseLogger",
  "levelname": "ERROR",
  "message": "Failed to add course. Missing fields: semester,
schedule, classroom, prerequisites, grading"
}
```

Setup and Execution

1. Clone the Repository:

```
git clone https://github.com/nishchaybhutoria/CS203_Lab_01-main.git
cd CS203_Lab_01-main
```

2. Set Up the Environment:

```
python -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate
pip install -r requirements.txt
```

3. Run Jaeger:

```
sudo docker run -d --name jaeger \
  -e COLLECTOR_ZIPKIN_HTTP_PORT=9411 \
  -p 5775:5775/udp \
  -p 6831:6831/udp \
  -p 6832:6832/udp \
  -p 5778:5778 \
  -p 16686:16686 \
  -p 14268:14268 \
  -p 14250:14250 \
  -p 9411:9411 \
  jaegertracing/all-in-one:1.41 --log-level=debug
```

4. Run the Flask Application:

```
flask run
```

5. Access the Application:

- Flask App: <http://localhost:5000>
- Jaeger UI: <http://localhost:16686>